

TM 83
100

File read in SVD fails,
though I'll ignore that.

Code compiles and
runs: +20

Final Part 1: Report

Julian Lehrer

Question 1a.

Using the singular value reconstruction with the $k \in \{20, 40, 80, 160, 320, 640, 1280, 2560, 3355\}$ largest singular values in the Σ matrix, we obtain the following compressed images:




(a) $k = 20$




(b) $k = 40$



(c) $k = 80$




(a) $k = 160$




(b) $k = 320$




(c) $k = 640$



(a) $k = 1280$



(b) $k = 2560$




(c) $k = 3355$ (original image)

SVD 25
30

Reported singular values?
-4

With more values in the Σ matrix of the SVD set to zero ($k = 20$ being the fewest nonzero), the visual distortion is the highest. However, by just increasing $k = 80$, the visual distortion is quite minimal, and for $k > 80$ I cannot see any visual difference with a small png. Now, defining the average Frobenius error of an $m \times n$ matrix W by $\|W\|_F / (mn)$, we plot the errors of the difference between the original image and the k th SVD compression, i.e. $\|A - A_{k\sigma}\|_F / (mn)$.

What is μ ?



-1 should take 1260...

With $k = 20$, we obtain an error less than 10^{-3} , and for subsequent k (with less distortion), we clearly do as well.

Question b.1

The Gauss-Jacobi and Gauss-Seidel are both iterative methods for solving exact systems of equations of the form $Ax = b$, where A is an $m \times m$ matrix and b is an m -vector. For the former, we factorize $A = D + L + U$ where D contains the diagonal elements and zeros elsewhere, and L and U contain the strictly lower and upper diagonal elements of A , respectively. Then we simply apply the iteration $x^{k+1} = D^{-1} (b - (L + U)x^k)$.

Note that for the Gauss-Jacobi method, the values x_i^k obtained in the k th it-

Gauss
Methods
Report
20
25

eration remain completely unchanged until the $k + 1$ th iteration has been fully calculated. Instead, the Gauss-Seidel uses the values x_i^{k+1} as soon as they're known. That is, once we calculate x_j^k in the j th equation we use it in the x_{j+1}^k th equation. This difference allows us to only store one solution vector x , instead of having to store two for the current and previous.

For the Gauss-Jacobi method to converge, it is sufficient (but not necessary) that A must be strictly diagonally dominant. More generally, we must have that $\rho(D^{-1}(L + U)) < 1$, that is, the spectrum is less than 1. The Gauss-Seidel method is known to converge if A is symmetric positive-definite, or if A is strictly or irreducibly diagonally dominant.

In our case, we define a 10×10 matrix A such that $a_{ij} = 1$ for $i \neq j$ and $a_{ii} = D$ for $D \in \{2, 5, 10, 100, 1000\}$. Below, visualize the convergence for the two methods, where $b_i = i$ for $i = 1, \dots, 10$.

Gauss-Jacobi, Log 2-Norm Error




Figure 4: Convergence of Gauss-Jacobi, number of iterations vs log norm error.

Gauss-Seidel, Log 2-Norm Error




Figure 5: Convergence of Gauss-Seidel, number of iterations vs log norm error

For the Jacobi method, we can see that for $D = 2, 5$, the method diverges, whereas for the Gauss-Seidel method, these values converge (albeit quite slowly). For the Gauss-Seidel algorithm, since $a_{ii} = D$ and $a_{ij} = 1$, $i \neq j$, we have that $a_{ji} = a_{ij}$ for all i and j . Therefore, A is symmetric. Now, consider \hat{x}_i to be the i th element of x under the image of our matrix A . Then $\hat{x}_i = Dx_i + \sum_{j=1}^m_{j \neq i} x_j$. Therefore, $x_i \hat{x}_i = D(x_i)^2 + \sum_{j=1}^m_{j \neq i} x_j x_i$. For large enough D , this means A is positive-definite, so the algorithm converges.

Question b.2 The conjugate gradient algorithm numerically solves an exact system of equations $Ax = b$ where A is symmetric and positive-definite. Then, this is equivalent to minimizing the function $f(x) = \frac{1}{2}x^T Ax - x^T b$, since $\nabla f = Ax - b$ and solving $\nabla f = 0$ is equivalent to solving $Ax = b$. The conjugate gradient method therefore chooses n conjugate directions, taking a step size towards the global minimum of f . The algorithm begins by choosing the direction like the gradient descent algorithm, i.e. opposite the direction of greatest ascent, or the gradient vector itself. So, $p_0 = b - Ax_0$. We then successively choose directions p_k orthogonal to p_{k-1} , and search along our surface until we find the minima.

Here, we show the number of iterations convergence takes for the Gauss-Jacobi, Gauss-Seidel, and Conjugate-Gradient methods.

Finite steps
to converge?
-2

CG/BJ/SS
-5

CG/BJ/SS
18
25

Your plots don't
show this either..

↓

D	Conjugate Gradient	Gauss-Jacobi	Gauss-Seidel
2	2	Not reached	58
5	2	Not reached	16
10	2	11	10
100	2	1	5
1000	2	1	4

In general, the diagonal preconditioner is particularly efficient for diagonally dominant matrices. However, in the case of the conjugate gradient method, multiplying by the diagonal preconditioner given by $M = \text{diag}(\frac{1}{D}, \dots, \frac{1}{D})$ means scaling each row by $\frac{1}{D}$, which in the case of our matrix is really just scaling the entire matrix down by a constant since the diagonals are constant. Therefore, we aren't actually changing the system in any way, and therefore the preconditioning isn't helpful at all. In the case where A is a 10×10 matrix with $a_{ii} = i$, we have that the conjugate gradient algorithm converges in 10 iterations, and for A being 100×100 with $a_{ii} = i$, it actually converges on the first iteration! This means that the initial vector x given by $x = [1, \dots, 1]^T$ (how I initialized x_0) was a particularly good guess in this case.

→ That is a bit unusual.
I suppose we should
make a particular
starting guess in
the future...